

# How to Get Data — An Introduction into **quantmod**

January 17, 2014

## 1 The S&P 500 index

This vignette gives a brief introduction to obtaining data from the web by using the R package **quantmod**. As example data, the time series of the S&P 500 index is used. This data is also used in Carmona, page 5 ff.

First, we load the package **quantmod**:

```
R> library("quantmod")
```

**quantmod** provides a very suitable function for downloading financial data from the web. This function is called `getSymbols`. The first argument of this function is a character vector specifying the names of the symbols to be loaded and the second one specifies the environment where the object is created. The help page of this function (`?getSymbols`) provides more information. By default, objects are created in the workspace. Here, we use a separate environment which we call `sp500` to store the downloaded data. We first create the environment:

```
R> sp500 <- new.env()
```

We can then download the S&P 500 timeseries (symbol: `GSPC`) from 1960-01-04 to 2009-01-01 from *yahoo finance* via:

```
R> getSymbols("~GSPC", env = sp500, src = "yahoo",  
+           from = as.Date("1960-01-04"), to = as.Date("2009-01-01"))
```

```
[1] "GSPC"
```

The package **quantmod** works with a variety of sources. Current `src` methods available are: *yahoo*, *google*, *MySQL*, *FRED*, *csv*, *RData*, and *oanda*. For example, FRED (Federal Reserve Economic Data), is a database of 20,070 U.S. economic time series (see <http://research.stlouisfed.org/fred2/>).

There are several possibilities, to load the variable `GSPC` from the environment `sp500` to a variable in the global environment (also known as the workspace), e.g., via

```
R> GSPC <- sp500$GSPC  
R> GSPC1 <- get("GSPC", envir = sp500)  
R> GSPC2 <- with(sp500, GSPC)
```

The object `GSPC1` and `GSPC2` are identical to `GSPC` so we can remove them from the workspace with:

```
R> rm(GSPC1)  
R> rm(GSPC2)
```

The function `head` shows the first six rows of the data.

```
R> head(GSPC)
```

	GSPC.Open	GSPC.High	GSPC.Low	GSPC.Close	GSPC.Volume	GSPC.Adjusted
1960-01-04	59.91	59.91	59.91	59.91	3990000	59.91
1960-01-05	60.39	60.39	60.39	60.39	3710000	60.39
1960-01-06	60.13	60.13	60.13	60.13	3730000	60.13
1960-01-07	59.69	59.69	59.69	59.69	3310000	59.69
1960-01-08	59.50	59.50	59.50	59.50	3290000	59.50
1960-01-11	58.77	58.77	58.77	58.77	3470000	58.77

The data object is an “extensible time series” (xts) object:

```
R> class(GSPC)
```

```
[1] "xts" "zoo"
```

Here, it is a multivariate (irregular) time series with 12334 daily observations on 6 variables:

```
R> dim(GSPC)
```

```
[1] 12334      6
```

Such xts objects allow for conveniently selecting single time series using \$

```
R> head(GSPC$GSPC.Volume)
```

	GSPC.Volume
1960-01-04	3990000
1960-01-05	3710000
1960-01-06	3730000
1960-01-07	3310000
1960-01-08	3290000
1960-01-11	3470000

as well as very conveniently selecting observations according to their time stamp by using a character “row” index in the ISO 8601 date/time format ‘CCYY-MM-DD HH:MM:SS’, where more granular elements may be left out in which case all observations with time stamp “matching” the given one will be used. E.g., to get all observations in March 1970:

```
R> GSPC["1970-03"]
```

	GSPC.Open	GSPC.High	GSPC.Low	GSPC.Close	GSPC.Volume	GSPC.Adjusted
1970-03-02	89.50	90.80	88.92	89.71	12270000	89.71
1970-03-03	89.71	90.67	88.96	90.23	11700000	90.23
1970-03-04	90.23	91.05	89.32	90.04	11850000	90.04
1970-03-05	90.04	90.99	89.38	90.00	11370000	90.00
1970-03-06	90.00	90.36	88.84	89.44	10980000	89.44
1970-03-09	89.43	89.43	87.94	88.51	9760000	88.51
1970-03-10	88.51	89.41	87.89	88.75	9450000	88.75
1970-03-11	88.75	89.58	88.11	88.69	9180000	88.69
1970-03-12	88.69	89.09	87.68	88.33	9140000	88.33
1970-03-13	88.33	89.43	87.29	87.86	9560000	87.86
1970-03-16	87.86	87.97	86.39	86.91	8910000	86.91

1970-03-17	86.91	87.86	86.36	87.29	9090000	87.29
1970-03-18	87.29	88.28	86.93	87.54	9790000	87.54
1970-03-19	87.54	88.20	86.88	87.42	8930000	87.42
1970-03-20	87.42	87.77	86.43	87.06	7910000	87.06
1970-03-23	87.06	87.64	86.19	86.99	7330000	86.99
1970-03-24	86.99	88.43	86.90	87.98	8840000	87.98
1970-03-25	88.11	91.07	88.11	89.77	17500000	89.77
1970-03-26	89.77	90.65	89.18	89.92	11350000	89.92
1970-03-30	89.92	90.41	88.91	89.63	9600000	89.63
1970-03-31	89.63	90.17	88.85	89.63	8370000	89.63

It is also possible to specify a range of timestamps using ‘/’ as the range separator, where both endpoints are optional: e.g.,

```
R> GSPC["/1960-01-06"]
```

	GSPC.Open	GSPC.High	GSPC.Low	GSPC.Close	GSPC.Volume	GSPC.Adjusted
1960-01-04	59.91	59.91	59.91	59.91	3990000	59.91
1960-01-05	60.39	60.39	60.39	60.39	3710000	60.39
1960-01-06	60.13	60.13	60.13	60.13	3730000	60.13

gives all observations up to Epiphany (Jan 6) in 1960, and

```
R> GSPC["2008-12-25/"]
```

	GSPC.Open	GSPC.High	GSPC.Low	GSPC.Close	GSPC.Volume	GSPC.Adjusted
2008-12-26	869.51	873.74	866.52	872.80	1880050000	872.80
2008-12-29	872.37	873.70	857.07	869.42	3323430000	869.42
2008-12-30	870.58	891.12	870.58	890.64	3627800000	890.64
2008-12-31	890.59	910.32	889.67	903.25	4172940000	903.25

gives all observations from Christmas (Dec 25) in 2008 onwards.

If we are interested in the daily values of the weekly last-traded-day, we aggregate it by using an appropriate function from the “zoo Quick-Reference” (Shah et al., 2005). The “zoo Quick-Reference” can be found in the web, [cran.r-project.org/web/packages/zoo/vignettes/zoo-quickref.pdf](http://cran.r-project.org/web/packages/zoo/vignettes/zoo-quickref.pdf), and it is strongly recommended to have a look at this vignette since it gives a very good overview of the **zoo** package. Their convenience function `nextfri` computes for each “Date” the next Friday.

```
R> nextfri <- function(x) 7 * ceiling(as.numeric(x - 5 + 4)/7) + as.Date(5 - 4)
```

We get the aggregated data then via

```
R> SP.we <- aggregate(GSPC, nextfri, tail, 1)
```

The function `aggregate` splits the data into subsets — here according to the function `nextfri` — and computes statistics for each, i.e., takes the last value, which is done by `tail`.

(This works because the data object is also a “Z’s ordered observations” (zoo) object which knows to apply `nextfri()` to the index (timestamps). Package **xts** also provides functionality for aggregating xts objects; e.g., `apply.weekly()`.)

By extracting the column `GSPC.Close` we derive the closing values for each last-traded-day

```
R> SPC.we <- SP.we$GSPC.Close
```

and create a plot of this time series via

```
R> plot(SPC.we)
```

(see Figure 1)




Figure 1: Plot of the weekly S&P 500 index closing values from 1960-01-04 to 2009-01-01.

## 2 Investigating the NASDAQ100 index

In this example we want to analyze an American stock exchange, the National Association of Securities Dealers Automated Quotations, better known as NASDAQ. It is the largest electronic screen-based equity securities trading market in the United States.

From the web we obtained a txt-file containing the names and symbols of the NASDAQ100 (June 2009) which can be loaded via

```
R> sym <- readLines("nasdaqsymbols.txt")

R> head(sym)

[1] "CTIVISION INC \tATVI"          "ADOBE SYS INC \tADBE"
[3] "ALTERA CORP \tALTR"           "AMAZON \tAMZN"
[5] "AMERICAN POWER CONVER CORP \tAPCC" "AMGEN \tAMGN"
```

The elements are separated by a space and a tab (`\t`). We therefore use the command `strsplit` to split the vector into names and symbols:

```
R> tmp <- strsplit(sym, split = " \t")
```

We can then use the function `lapply` to get only the symbol of each company, and `unlist` to create a character vector with the symbols of the NASDAQ companies as entries.

```
R> symbols <- unlist(lapply(tmp, `[`, 2))
```

We assign the companies' names to the symbols

```
R> names(symbols) <- unlist(lapply(tmp, `[`, 1))
```

and get

```
R> head(symbols, 10)
```

CTIVISION INC	ADOBE SYS INC
"ATVI"	"ADBE"
ALTERA CORP	AMAZON
"ALTR"	"AMZN"
AMERICAN POWER CONVER CORP	AMGEN
"APCC"	"AMGN"
APOLLO GROUP-A	APPLE COMPUTER
"APOL"	"AAPL"
APPLIED MATERIAL	ATI TECHNOLOGIES
"AMAT"	"ATYT"

As before we create a new environment for our NASDAQ data and use the function `getSymbols` of the **quantmod** package to download the NASDAQ time series from 2000-01-01 to today. By using the command `tryCatch` we handle unusual conditions, including errors and warnings. In this case, if the data from a company are not available from yahoo finance, the message "Symbol ... not downloadable!" is provided.

```
R> nasdaq <- new.env()
R> for(i in head(symbols, 10)){
+   status <- tryCatch(getSymbols(i, env = nasdaq, src = "yahoo",
+                               from = as.Date("2000-01-01")),
+                     error = identity)
+   if(inherits(status, "error"))
+     cat("Symbol '", i, "' not downloadable!\n", sep = "")
+ }
```

Symbol 'ATYT' not downloadable!

E.g., the first values of the Apple NASDAQ time series are

```
R> with(nasdaq, head(AAPL))
```

	AAPL.Open	AAPL.High	AAPL.Low	AAPL.Close	AAPL.Volume	AAPL.Adjusted
2000-01-03	104.87	112.50	101.69	111.94	19144400	27.22
2000-01-04	108.25	110.62	101.19	102.50	18310000	24.92
2000-01-05	103.75	110.56	103.00	104.00	27818000	25.29
2000-01-06	106.12	107.00	95.00	95.00	27443200	23.10
2000-01-07	96.50	101.00	95.50	99.50	16463200	24.19
2000-01-10	102.00	102.25	94.75	97.75	18059200	23.77

Further, the command `chartSeries` of the package **quantmod** provides the full financial charting abilities to R and allows for an interaction within the charts. E.g., using

```
R> chartSeries(nasdaq$AAPL)
```

gives a chart of the Apple values (see Figure 2) and e.g., with the command `with(nasdaq, addOBV(AAPL))` the On-Balance volume can be visualized in the plot. See the manual of the **quantmod** package (Ryan, 2008) for the whole list of available plot and visualization functions.




Figure 2: barChart of Apple.

E.g., Bollinger bands consist of a center line and two price channels (bands) above and below it. The center line is an exponential moving average; the price channels are the standard deviations of the stock being studied. The bands will expand and contract as the price action of an issue becomes volatile (expansion) or becomes bound into a tight trading pattern (contraction).

We can add the Bollinger Bands to a plot by using the command: `addBBands(n = 20, sd = 2, ma = "SMA", draw = "bands", on = -1)`, where `n` denotes the number of moving average periods, `sd` the number of standard deviations and `ma` the used moving average process.

Have a look at the **quantmod** homepage for further examples and try to reproduce them, <http://www.quantmod.com/examples/intro/>.